Performance Attribution for Equity Portfolios

by Yang Lu and David Kane

Abstract

Almost all portfolio managers measure performance with reference to a benchmark. The difference in return between a portfolio and the benchmark is its active return. Performance attribution decomposes the active return. The two most common approaches are the Brinson model and a regression-based analysis.[a]

The **pa** package provides tools for conducting both methods for equity portfolios. The Brinson model takes an ANOVA-type approach and decomposes the active return of any portfolio into asset allocation, stock selection, and interaction effects. The regression-based analysis utilizes estimated coefficients from a linear model to estimate the contributions from different factors.

---

[a]See Brinson et al. (1986) and Grinold (2006) for more information.

Introduction

Almost all portfolio managers measure performance with reference to a benchmark. The difference in return between a portfolio and the benchmark is its active return. Performance attribution decomposes the active return. The two most common approaches are the Brinson model and a regression-based analysis.[1]

The **pa** package provides tools for conducting both methods for equity portfolios. The Brinson model takes an ANOVA-type approach and decomposes the active return of any portfolio into asset allocation, stock selection, and interaction effects. The regression-based analysis utilizes estimated coefficients from a linear model to estimate the contributions from different factors.

Data

We demonstrate the use of the **pa** package with a series of examples based on data from MSCI Barra's Global Equity Model II (GEM2).[2] The original data set contains selected attributes such as industry, size, country, and various style factors for a universe of approximately 48,000 securities on a monthly basis. For illustrative purposes, this article uses three modified versions of the original data set (`year`, `quarter`, and `jan`), each containing 3000 securities. The data frame, `quarter`, is a subset of `year`, containing the data of the first quarter. The data frame, `jan`, is a subset of `quarter` with the data from January, 2010.

```
> data(year)
> names(year)

 [1] "barrid"    "name"      "return"
 [4] "date"      "sector"    "momentum"
 [7] "value"     "size"      "growth"
[10] "cap.usd"   "yield"     "country"
[13] "currency"  "portfolio" "benchmark"
```

See `?year` for information on the different variables. The top 200 securities, based on `value` scores, in January are selected as portfolio holdings and are held through December 2010 with monthly rebalances to maintain equal-weighting. The benchmark for this portfolio is defined as the largest 1000 securities based on `size` each month. The benchmark is cap-weighted.

Here is a sample of rows and columns from the data frame `year`:

```
                                     name
44557  BLUE STAR OPPORTUNITIES CORP
25345  SEADRILL
264017 BUXLY PAINTS (PKR10)
380927 CDN IMPERIAL BK OF COMMERCE
388340 CDN IMPERIAL BK OF COMMERCE
          return        date      sector   size
44557    0.00000  2010-01-01      Energy   0.00
25345   -0.07905  2010-01-01      Energy  -0.27
264017  -0.01754  2010-05-01   Materials   0.00
380927   0.02613  2010-08-01  Financials   0.52
388340  -0.00079  2010-11-01  Financials   0.55
        country portfolio benchmark
44557       USA     0.000  0.000000
25345       NOR     0.000  0.000427
264017      PAK     0.005  0.000000
380927      CAN     0.005  0.000012
388340      CAN     0.005  0.000012
```

The portfolio has 200 equal-weighted holdings each month. The row for Canadian Imperial Bank of Commerce indicates that it is one of the 200 portfolio holdings with a weight of 0.5% in 2010. Its return was 2.61% in August, and close to flat in November.

The Brinson Model

Consider an equity portfolio manager who uses the S&P 500 as the benchmark. In a given month, she outperformed the S&P by 3%. Part of that performance was due to the fact that she allocated more

---

[1]See Brinson et al. (1986) and Grinold (2006) for more information.
[2]See www.msci.com and Menchero et al. (2008) for more information.

weight of the portfolio to certain sectors that performed well. Call this the *allocation effect*. Part of her outperformance was due to the fact that some of the stocks she selected did better than their sector as a whole. Call this the *selection effect*. The residual can then be attributed to an interaction between allocation and selection – the *interaction effect*. The Brinson model provides mathematical definitions for these terms and methods for calculating them.

The example above uses sector as the classification scheme when calculating the allocation effect. But the same approach can work with any other variable which places each security into one, and only one, discrete category: country, industry, and so on. In fact, a similar approach can work with continuous variables that are split into discrete ranges: the highest quintile of market cap, the second highest quintile and so forth. For generality, we will use the term "category" to describe any classification scheme which places each security in one, and only one, category.

Notations:

- w_i^B is the weight of security i in the benchmark.

- w_i^P is the weight of security i in the portfolio.

- W_j^B is the weight of category j in the benchmark. $W_j^B = \sum w_i^B$, $i \in j$.

- W_j^P is the weight of a category j in the portfolio. $W_j^P = \sum w_i^P$, $i \in j$.

- The sum of the weight w_i^B, w_i^P, W_j^B, and W_j^P is 1, respectively.

- r_i is the return of security i.

- R_j^B is the return of a category j in the benchmark. $R_j^B = \sum w_i^B r_i$, $i \in j$.

- R_j^P is the return of a category j in the portfolio. $R_j^P = \sum w_i^P r_i$, $i \in j$.

The return of a portfolio, R_P, can be calculated in two ways:

- On an individual security level by summing over n stocks: $R_P = \sum\limits_{i=1}^{n} w_i^P r_i$.

- On a category level by summing over N categories: $R_P = \sum\limits_{j=1}^{N} W_j^P R_j^P$.

Similar definitions apply to the return of the benchmark, R_B,

- $R_B = \sum\limits_{i=1}^{n} w_i^B r_i$.

- $R_B = \sum\limits_{j=1}^{N} W_j^B R_j^B$.

Active return of a portfolio, R_{active}, is a performance measure of a portfolio relative to its benchmark. The two conventional measures of active return are arithmetic and geometric. The **pa** package implements the arithmetic measure of the active return for a single-period Brinson model because an arithmetic difference is more intuitive than a ratio over a single period.

The arithmetic active return of a portfolio, R_{active}, is the portfolio return R_P less the benchmark return R_B:

$$R_{active} = R_P - R_B.$$

Since the category weights of the portfolio are generally different from those of the benchmark, allocation plays a role in the active return, R_{active}. The same applies to stock selection effects. Within a given category, the portfolio and the benchmark will rarely have exactly the same holdings. Allocation effect $R_{allocation}$ and selection effect $R_{selection}$ over N categories are defined as:

$$R_{allocation} = \sum\limits_{j=1}^{N} W_j^P R_j^B - \sum\limits_{j=1}^{N} W_j^B R_j^B,$$

$$R_{selection} = \sum\limits_{j=1}^{N} W_j^B R_j^P - \sum\limits_{j=1}^{N} W_j^B R_j^B.$$

The intuition behind the allocation effect is that a portfolio would produce different returns with different allocation schemes (W_j^P vs. W_j^B) while having the same stock selection and thus the same return (R_j^B) for each category. The difference between the two returns, caused by the allocation scheme, is called the allocation effect ($R_{allocation}$). Similarly, two different returns can be produced when two portfolios have the same allocation (W_j^B) yet dissimilar returns due to differences in stock selection within each category (R_j^p vs. R_j^B). This difference is the selection effect ($R_{selection}$).

Interaction effect, $R_{interaction}$, is the result of subtracting return due to allocation $R_{allocation}$ and return due to selection $R_{selection}$ from the active return R_{active}:

$$R_{interaction} = R_{active} - R_{allocation} - R_{selection}.$$

The Brinson model allows portfolio managers to analyze the active return of a portfolio using any attribute of a security, such as country or sector. Unfortunately, it is very hard to expand the analysis beyond two categories. As the number of categories increases, this procedure is subject to the curse of dimensionality. To some extent, the regression-based model detailed later ameliorates this problem.

Brinson Tools

Brinson analysis is run by calling the function `brinson` to produce an object of class `brinson`.

```
> data(jan)
> br.single <- brinson(x = jan, date.var = "date",
+               cat.var = "sector",
+               bench.weight = "benchmark",
+               portfolio.weight = "portfolio",
+               ret.var = "return")
```

The data frame, `jan`, contains all the information necessary to conduct a single-period Brinson analysis. `date.var`, `cat.var`, and `return` identify the columns containing the date, the factor to be analyzed, and the return variable, respectively. `bench.weight` and `portfolio.weight` specify the name of the benchmark weight column and that of the portfolio weight column in the data frame.

Calling `summary` on the resulting object `br.single` of class `brinson` reports essential information about the input portfolio (including the number of securities in the portfolio and the benchmark as well as sector exposures) and the results of the Brinson analysis.

```
> summary(br.single)
```

```
Period:                          2010-01-01
Methodology:                     Brinson
Securities in the portfolio:     200
Securities in the benchmark:     1000
```

```
Exposures
            Portfolio Benchmark
Energy          0.085    0.2782
Materials       0.070    0.0277
Industrials     0.045    0.0330
ConDiscre       0.050    0.0188
ConStaples      0.030    0.0148
HealthCare      0.015    0.0608
Financials      0.370    0.2979
InfoTech        0.005    0.0129
TeleSvcs        0.300    0.1921
Utilities       0.030    0.0640
```

```
Returns
                     2010-01-01
Allocation Effect      -0.00140
Selection Effect        0.01418
Interaction Effect      0.00191
Active Return           0.01469
```

The `br.single` summary shows that the active return of the portfolio, in January, 2010 was 1.47%. This return can be decomposed into allocation effect (-0.14%), selection effect (1.42%), and interaction effect (0.19%).

```
> plot(br.single, var = "sector", type = "return")
```


Figure 1: Sector Return.

Figure 1 is a visual representation of the return of both the portfolio and the benchmark sector by sector in January, 2010. Utilities was the sector with the highest active return in the portfolio.

To obtain Brinson attribution on a multi-period data set, one calculates allocation, selection and interaction within each period and aggregates them across time. There are three methods for this – arithmetic, geometric, and optimized linking. The arithmetic attribution model calculates active return and contributions due to allocation, selection, and interaction in each period and sums them over multiple periods.

In practice, analyzing a single-period portfolio is meaningless as portfolio managers and their clients are more interested in the performance of a portfolio over multiple periods. To apply the Brinson model over time, we can use the function `brinson` and input a multi-period data set (for instance, `quarter`) as shown below.

```
> data(quarter)
> br.multi <- brinson(quarter, date.var = "date",
+               cat.var = "sector",
+               bench.weight = "benchmark",
+               portfolio.weight = "portfolio",
+               ret.var = "return")
```

The object `br.multi` of class `brinsonMulti` is an example of a multi-period Brinson analysis.

```
> exposure(br.multi, var = "size")
```

```
$Portfolio
     2010-01-01 2010-02-01 2010-03-01
Low       0.140      0.140      0.155
2         0.050      0.070      0.045
3         0.175      0.145      0.155
4         0.235      0.245      0.240
High      0.400      0.400      0.405
```

```
$Benchmark
      2010-01-01 2010-02-01 2010-03-01
Low       0.0681     0.0568     0.0628
2         0.0122     0.0225     0.0170
3         0.1260     0.1375     0.1140
4         0.2520     0.2457     0.2506
High      0.5417     0.5374     0.5557
```

The `exposure` method on the class `br.multi` object shows the exposure of the portfolio and the benchmark based on a user-specified variable. Here, it shows the exposure on `size`. We can see that the portfolio overweights the benchmark in the lowest quintile in `size` and underweights in the highest quintile.

```
> returns(br.multi, type = "arithmetic")

$Raw
                2010-01-01 2010-02-01 2010-03-01
Allocation         -0.0014     0.0062     0.0047
Selection           0.0142     0.0173    -0.0154
Interaction         0.0019    -0.0072    -0.0089
Active Return       0.0147     0.0163    -0.0196

$Aggregate
                2010-01-01, 2010-03-01
Allocation                      0.0095
Selection                       0.0160
Interaction                    -0.0142
Active Return                   0.0114
```

The `returns` method shows the results of the Brinson analysis applied to the data from January, 2010 through March, 2010. The first portion of the `returns` output shows the Brinson attribution in individual periods. The second portion shows the aggregate attribution results. The portfolio formed by top 200 value securities in January had an active return of 1.14% over the first quarter of 2010. The allocation and the selection effects contributed 0.95% and 1.6% respectively; the interaction effect decreased returns by 1.42%.

Regression

One advantage of a regression-based approach is that such analysis allows one to define their own attribution model by easily incorporating multiple variables in the regression formula. These variables can be either discrete or continuous.

Suppose a portfolio manager wants to find out how much each of the value, growth, and momentum scores of her holdings contributes to the overall performance of the portfolio. Consider the following linear regression without the intercept term based on a single-period portfolio of n securities with k different variables:

$$\mathbf{r}_n = \mathbf{X}_{n,k}\mathbf{f}_k + \mathbf{u}_n$$

where

- $\mathbf{r}_n$ is a column vector of length n. Each element in $\mathbf{r}_n$ represents the return of a security in the portfolio.

- $\mathbf{X}_{n,k}$ is an n by k matrix. Each row represents k attributes of a security. There are n securities in the portfolio.

- $\mathbf{f}_k$ is a column vector of length k. The elements are the estimated coefficients from the regression. Each element represents the *factor return* of an attribute.

- $\mathbf{u}_n$ is a column vector of length n with residuals from the regression.

In the case of this portfolio manager, suppose that she only has three holdings in her portfolio. r_3 is thus a 3 by 1 matrix with returns of all her three holdings. The matrix $\mathbf{X}_{3,3}$ records the score for each of the three factors (value, growth, and momentum) in each row. $\mathbf{f}_3$ contains the estimated coefficients of a regression $\mathbf{r}_3$ on $\mathbf{X}_{3,3}$.

The active exposure of each of the k variables, X_i, $i \in k$, is expressed as

$$X_i = \mathbf{w}_{active}'\mathbf{x}_{n,i},$$

where X_i is the value representing the active exposure of the attribute i in the portfolio, $\mathbf{w}_{active}$ is a column vector of length n containing the active weight of every security in the portfolio, and $\mathbf{x}_{n,i}$ is a column vector of length n with attribute i for all securities in the portfolio. Active weight of a security is defined as the difference between the portfolio weight of the security and its benchmark weight.

Using the example mentioned above, the active exposure of the attribute `value`, X_{value} is the product of $\mathbf{w}_{active}'$ (containing active weight of each of the three holdings) and $\mathbf{x}_3$ (containing value scores of the three holdings).

The contribution of a variable i, R_i, is thus the product of the factor returns for the variable i, f_i and the active exposure of the variable i, X_i. That is,

$$R_i = f_i X_i.$$

Continuing the example, the contribution of value is the product of f_{value} (the estimated coefficient for value from the linear regression) and X_{value} (the active exposure of value as shown above).

Therefore, the active return of the portfolio R_{active} is the sum of contributions of all k variables and the residual u (a.k.a. the interaction effect),

$$R_{active} = \sum_{i=1}^{k} R_i + u.$$

For instance, a hypothetical portfolio has three holdings (A, B, and C), each of which has two attributes – size and value.

```
  Return Name Size Value Active_Weight
1    0.3    A  1.2   3.0           0.5
2    0.4    B  2.0   2.0           0.1
3    0.5    C  0.8   1.5          -0.6
```

Following the procedure as mentioned, the factor returns for size and value are -0.0313 and -0.1250. The active exposure of size is 0.32 and that of value is 0.80. The active return of the portfolio is -11% which can be decomposed into the contribution of size and that of value based on the regression model. Size contributes 1% of the negative active return of the portfolio and value causes the portfolio to lose the other 10.0%.

Regression Tools

The **pa** package provides tools to analyze both single-period and multi-period data frames.

```
> rb.single <- regress(jan, date.var = "date",
+                      ret.var = "return",
+                      reg.var = c("sector", "growth",
+                       "size"),
+                      benchmark.weight = "benchmark",
+                      portfolio.weight = "portfolio")
> exposure(rb.single, var = "growth")

      Portfolio Benchmark
Low       0.305    0.2032
2         0.395    0.4225
3         0.095    0.1297
4         0.075    0.1664
High      0.130    0.0783
```

`reg.var` specifies the columns containing variables whose contributions are to be analyzed. Each of the `reg.var` input corresponds to a particular column in $\mathbf{X}_{n,k}$ from the aforementioned regression model. `ret.var` specifies the column in the data frame `jan` based on which $\mathbf{r}_n$ in the regression model is formed.

```
> exposure(rb.single, var = "growth")

      Portfolio Benchmark
Low       0.305    0.2032
2         0.395    0.4225
3         0.095    0.1297
4         0.075    0.1664
High      0.130    0.0783
```

Calling `exposure` with a specified `var` yields information on the exposure of both the portfolio and the benchmark by that variable. If `var` is a continuous variable, for instance, `growth`, the exposure will be shown in 5 quantiles. Majority of the high `value` securities in the portfolio in January have relatively low `growth` scores.

```
> summary(rb.single)
```

```
Period:                              2010-01-01
Methodology:                         Regression
Securities in the portfolio:         200
Securities in the benchmark:         1000

Returns
                    2010-01-01
Sector               0.003189
Growth               0.000504
Size                 0.002905
Residual             0.008092
Portfolio Return    -0.029064
Benchmark Return    -0.043753
Active Return        0.014689
```

The `summary` method shows the number of securities in the portfolio and the benchmark, and the contribution of each input variable according to the regression-based analysis. In this case, the portfolio made a loss of 2.91% and the benchmark lost 4.38%. Therefore, the portfolio outperformed the benchmark by 1.47%. `sector`, `growth`, and `size` contributed 0.32%, 0.05%, and 0.29%, respectively.

Regression-based analysis can be applied to a multi-period data frame by calling the same method `regress`. By typing the name of the class object `rb.multi` directly, a short summary of the analysis is provided, showing the starting and ending period of the analysis, the methodology, and the average number of securities in both the portfolio and the benchmark.

```
> rb.multi <- regress(year, date.var = "date",
+                     ret.var = "return",
+                     reg.var = c("sector", "growth",
+                      "size"),
+                     benchmark.weight = "benchmark",
+                     portfolio.weight = "portfolio")
> rb.multi

Period starts:                       2010-01-01
Period ends:                         2010-12-01
Methodology:                         Regression
Securities in the portfolio:         200
Securities in the benchmark:         1000
```

The regression-based summary shows that the contribution of each input variable in addition to the basic information on the portfolio. The summary suggests that the active return of the portfolio in year 2010 is 10.1%. The `Residual` number indicates the contribution of the interaction among various variables including `sector`, `growth`, and `growth`. Based on the regression model, `size` contributed to the lion share of the active return.

```
> summary(rb.multi)

Period starts:                       2010-01-01
Period ends:                         2010-12-01
Methodology:                         Regression
```

```
Avg securities in the portfolio:      200
Avg securities in the benchmark:     1000

Returns
$Raw
                 2010-01-01 2010-02-01 2010-03-01
Sector               0.0032     0.0031     0.0002
Growth               0.0005     0.0009    -0.0001
Size                 0.0029     0.0295     0.0105
Residual             0.0081    -0.0172    -0.0302
Portfolio Return    -0.0291     0.0192     0.0298
Benchmark Return    -0.0438     0.0029     0.0494
Active Return        0.0147     0.0163    -0.0196
                 2010-04-01 2010-05-01 2010-06-01
Sector               0.0016     0.0039     0.0070
Growth               0.0001     0.0002     0.0004
Size                 0.0135     0.0037     0.0018
Residual            -0.0040     0.0310     0.0183
Portfolio Return    -0.0080    -0.0381     0.0010
Benchmark Return    -0.0192    -0.0769    -0.0266
Active Return        0.0113     0.0388     0.0276
                 2010-07-01 2010-08-01 2010-09-01
Sector               0.0016     0.0047    -0.0022
Growth              -0.0005     0.0005    -0.0006
Size                 0.0064     0.0000     0.0096
Residual            -0.0324     0.0173    -0.0220
Portfolio Return     0.0515    -0.0119     0.0393
Benchmark Return     0.0764    -0.0344     0.0545
Active Return       -0.0249     0.0225    -0.0152
                 2010-10-01 2010-11-01 2010-12-01
Sector               0.0015    -0.0044    -0.0082
Growth              -0.0010    -0.0004     0.0010
Size                 0.0022     0.0130     0.0056
Residual             0.0137     0.0175    -0.0247
Portfolio Return     0.0414    -0.0036     0.0260
Benchmark Return     0.0249    -0.0293     0.0525
Active Return        0.0165     0.0257    -0.0265

$Aggregate
                 2010-01-01, 2010-12-01
Sector                           0.0120
Growth                           0.0011
Size                             0.1030
Residual                        -0.0269
Portfolio Return                 0.1191
Benchmark Return                 0.0176
Active Return                    0.1015
```

Figure 2 displays both the cumulative portfolio and benchmark returns from January, 2010 through December, 2010. It suggests that the portfolio, consisted of high `value` securities in January, consistently outperformed the benchmark in 2010. Outperformance in May and June helped the overall positive active return in 2010 to a large extent.

```
> plot(rb.multi, var = "sector", type = "return")
```
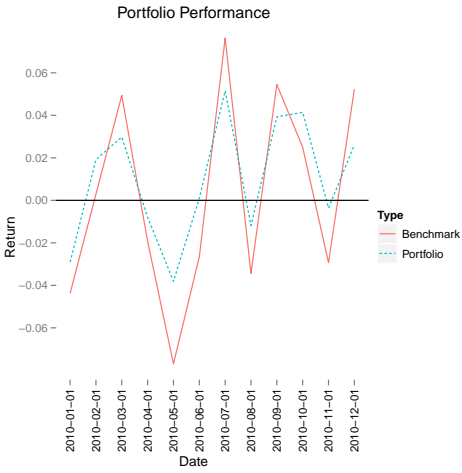

Figure 2: Performance Attribution.

Conclusion

In this paper, we describe two widely-used methods for performance attribution – the Brinson model and the regression-based approach, and provide a simple collection of tools to implement these two methods in R with the `pa` package. A comprehensive package, `portfolio` Enos and Kane (2006), provides facilities to calculate exposures and returns for equity portfolios. It is possible to use the `pa` package based on the output from the `portfolio` package. Further, the flexibility of R itself allows users to extend and modify these packages to suit their own needs and/or execute their preferred attribution methodology. Before reaching that level of complexity, however, `pa` provides a good starting point for basic performance attribution.

Bibliography

G. Brinson, R. Hood, and G. Beebower. Determinants of Portfolio Performance. *Financial Analysts Journal*, 42(4):39–44, Jul.–Aug. 1986. URL http://www.jstor.org/stable/4478947.

J. Enos and D. Kane. Analysing Equity Portfolios in R. *R News*, 6(2):13–19, MAY 2006. URL http://CRAN.R-project.org/doc/Rnews.

R. Grinold. Attribution, Modeling asset characteristics as portfolios. *The Journal of Portfolio Management*, page 19, Winter 2006.

J. Menchero, A. Morozov, and P. Shepard. The Barra Global Equity Model (GEM2). *MSCI Barra Research Notes*, Sept. 2008.

Yang Lu
Williams College

2896 Paresky Center, Williamstown, MA 01267
United States
yang.lu@williams.edu

David Kane
Affiliation
Address
United States
dave.kane@gmail.com